

M3-L1-P3

September 23, 2023

1 Problem 3 (5 points)

```
[ ]: import numpy as np
import matplotlib.pyplot as plt

def plot_data(data, c, title="",
    ↳ xlabel="$x_1$", ylabel="$x_2$", classes=["", ""], alpha=1):
    N = len(c)
    colors = ['royalblue', 'crimson']
    symbols = ['o', 's']

    plt.figure(figsize=(5,5), dpi=120)

    for i in range(2):
        x = data[:,0][c==i]
        y = data[:,1][c==i]

        plt.
        ↳ scatter(x,y,color=colors[i],marker=symbols[i],edgecolor="black",linewidths=0.
        ↳ 4,label=classes[i],alpha=alpha)

    plt.legend(loc="upper right")
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    ax = plt.gca()
    ax.set_xticklabels([])
    ax.set_yticklabels([])
    plt.xlim([-0.05,1.05])
    plt.ylim([-0.05,1.05])
    plt.title(title)

def plot_contour(predict, mapXY = None):
    res = 500
    vals = np.linspace(-0.05,1.05,res)
    x,y = np.meshgrid(vals,vals)
    XY = np.concatenate((x.reshape(-1,1),y.reshape(-1,1)),axis=1)
```

```

if mapXY is not None:
    XY = mapXY(XY)
contour = predict(XY).reshape(res, res)
plt.contour(x, y, contour)

```

1.1 Generate Dataset

(Don't edit this code.)

```


[ ]: def sample_ring(N,x,y,ro,ri):
    theta = np.random.rand(N)*2*np.pi
    r = np.random.rand(N)
    r = np.sqrt(r*(ro**2-ri**2)+ri**2)
    xs = x + r * np.cos(theta)
    ys = y + r * np.sin(theta)
    return xs, ys

def get_ring_dataset():
    np.random.seed(0)
    c0 = sample_ring(70,0.5,0.5,0.5,0.3)
    c1 = sample_ring(60,0.45,0.47,0.36,0.15)
    xs = np.concatenate([c0[0],c1[0]],0)
    ys = np.concatenate([c0[1],c1[1]],0)
    c = np.concatenate([np.zeros(70),np.ones(60)],0)
    return np.vstack([xs,ys]).T, c

[ ]: data, classes = get_ring_dataset()
format = dict(xlabel="Initial Speed",ylabel="Launch Angle", classes=["0 - Misses target", "1 - Hits target"])

plot_data(data, classes, **format)

```



1.2 Feature Expansion

Define a function to expand 2 features into more features For the features x_1 and x_2 , expand into:

- 1 - $x_1 - x_2 - x_1^2 - x_2^2 - \sin(x_1) - \cos(x_1) - \sin(x_2) - \cos(x_2) - \sin^2(x_1) - \cos^2(x_1) - \sin^2(x_2) - \cos^2(x_2)$
- $\exp(x_1) - \exp(x_2)$

```
[ ]: def feature_expand(x):
    x1 = x[:,0].reshape(-1, 1)
    x2 = x[:,1].reshape(-1, 1)

    # YOUR CODE GOES HERE:
    columns = [np.ones_like(x1), x1, x2, x1**2, x2**2,
               np.sin(x1), np.cos(x1), np.sin(x2), np.cos(x2),
               np.sin(x1)**2, np.cos(x1)**2, np.sin(x2)**2, np.cos(x2)**2,
               np.exp(x1), np.exp(x2)] # Add all expanded features to this list

    X = np.concatenate(columns, axis=1)
    return X

features = feature_expand(data)
print("Dataset size:", np.shape(data))
```

```
print("Expanded dataset size:", np.shape(features))
```

Dataset size: (130, 2)

Expanded dataset size: (130, 15)

1.3 Logistic Regression

Use SciKit-Learn's Logistic Regression model to learn the decision boundary for this data, using regularization. (The `C` argument controls regularization strength.)

Train this model on your expanded feature set.

Details about how to use this are here: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html


Notes: - λ is related to sklearn's regularization strength C by: $\lambda = 1/C$ - You may want to increase the maximum number of iterations


```
[ ]: from sklearn.linear_model import LogisticRegression


def get_logistic_regressor(features, classes, L = 1):
    # YOUR CODE GOES HERE
    # - Instantiate model with regularization
    # - Fit model to expanded data
    model = LogisticRegression(C=1/L, max_iter=10000)
    model.fit(features, classes)

    return model

[ ]: for L in [1e-9, 1e-1, 1]:
    model = get_logistic_regressor(features, classes, L)
    plot_data(data, classes, **format, title=f"$\lambda={L}$")
    plot_contour(model.predict, feature_expand)
    plt.show()
```







As λ increases, note what happens to the decision boundary. Why does this occur?

Lambda smooths out the fitted boundary. This results in not over fitting but if taken too far, as with $\lambda = 1$, the boundary doesn't fit the data well at all.